



## CONTACT

Phone · +31 6 16 95 82 82

Email · [abdurakhmanovtimur472@gmail.com](mailto:abdurakhmanovtimur472@gmail.com)

[LinkedIn · in/tymur-abdurakhmanov](https://www.linkedin.com/in/tymur-abdurakhmanov)

[@diklinuks](https://github.com/diklinuks)

Full portfolio with live demos of my projects:

[diklinuks.github.io/CV-WEB/](https://diklinuks.github.io/CV-WEB/)

## DETAILS

Born 29 Nov 2004

Available Sep 2026 – Feb 2027

Eindhoven, NL

Willing to relocate within NL

Nuffic-eligible

## LANGUAGES

English — full professional

Ukrainian — native

Russian — native

## EDUCATION

### HBO-ICT, BSc

Fontys University of Applied Sciences, Eindhoven · 2024–2028

Software Engineering, then Machine Learning, then Generative AI.

# Tymur Abdurakhmanov

Generative AI / LLM Engineer Intern · also open to Machine Learning Engineer

## PROFILE

Second-year HBO-ICT student at Fontys, focused on generative AI / LLM agents and machine learning. I put as much effort into the research side of a project as into the building: reading into a problem, comparing options, and writing down why I chose what I did. I pick up new tools quickly, I'm easy to get along with and communicate clearly, and I'm comfortable leading a team and taking responsibility for my work.

## SKILLS

### Generative AI & LLM agents

LLMs, RAG, AI agents, Multi-agent systems, LangChain, LangGraph, CrewAI, MCP, n8n, Prompt engineering, Gemini, GPT, Claude, Grok, pgvector, ChromaDB

### Machine Learning & Computer Vision

PyTorch, scikit-learn, YOLOv8, OpenCV, Object detection, Explainable AI (SHAP), pandas, NumPy

### Engineering & Data

Python, JavaScript, FastAPI, PostgreSQL, SQL, Docker, Git, Linux, Streamlit, REST APIs

## INDUSTRY & CLIENT PROJECTS

### Laméco — Generative AI compliance tool

GenAI semester · team of two · 2026

Built a tool that automatically checks a website's cookie use against GDPR and flags the violations. Reads an existing cookie report or crawls the site itself, catching cookies that load before the visitor consents. Generates a written analysis and emails it back when a report is sent in. Specifics under confidentiality.

### Brunel — Recruitment-campaign analytics

ML semester · team project · 2025–2026

Brunel is an international recruitment company. Built an application that, for each job title, analyses the job-advertising campaign data and recommends which advertising platform to use and why. Compared several model approaches and analysed how the different platforms performed across their campaigns. Specifics under confidentiality.

## PERSONAL PROJECTS

### Betsy — Autonomous procurement agent

Generative AI · live demo

As much research as code: did the domain and technical research, then wrote the architecture, decision logs, design and test scenarios before building. Betsy is an AI agent that runs purchasing for a small manufacturer across a **90-day** simulation. Monitors stock and reorders before items hit their reorder threshold, preventing stockouts. Scores and ranks suppliers by performance, then places orders with the best ones and tracks KPIs throughout. Learns from the orders a human rejects and from sudden events, and pauses for human approval on the bigger decisions. Every rule is explained in the documentation and enforced in the code.

### Eva — Multi-agent marketing team

Generative AI · live demo

Give it something to market (a product, a book, a service) and a team of specialised AI agents collaborates to produce the whole campaign audience, strategy, copy and creative concepts instead of one model doing everything. Currently building and testing how the agents split the work and hand off to each other.

### Floating Waste Detection

Computer vision · live demo

Detects floating waste (bottles, cans, cartons, plastic, paper) in water; runs live in the browser and scores video frame by frame. Built a custom dataset of **~9,000** images. Trained and compared **seven** models (from Random Forest and SVM up to YOLO) and selected the best (**mAP@50 ≈ 0.87**).

### Amsterdam Airbnb — price & market analysis

ML + data analysis · live demo

Set out to predict nightly prices for **~7,800** Amsterdam listings; the data couldn't support it (best model **R² ≈ 0.23**). Used explainability tools (SHAP) to show why the features fall short. Switched to clustering the listings into **five** market segments, shown on an interactive map.